

Chapter 1

Configuring Logging

1.1 Understanding rsyslogd and journald logging

On RHEL 7 there are two systems responsible for logging : **rsyslogd** and **journald**. These two services together to handle logging system information.

It is up to the services (containing the logging information) to decide how and where the log files will be written to. It is possible to write directly to a log file anywhere (e.g., /somewhere/my.log). The service may also choose to pass over the information to **systemctl** as well. The **systemctl** utility is used to start the service and keep track of the actions of the service while it's starting. Anything that goes through **systemd** will be writing to **journald**, which is the **systemd** way of logging.




Figure 1.1: Logging Options

The classic way of logging is through the **rsyslogd** daemon, which typically writes information to the **/var/log/** directory. Alternative locations can also be used.

Given the various ways to log information, it may also become challenging for a user to get that logged information. It can be through **journalctl** or **rsyslog**. It is possible to tie these two systems together.

1.1.1 Sharing logging information

To ensure that journalctl information is automatically logged to rsyslog, we need to add a couple of lines to `/etc/rsyslog.conf` and `/etc/rsyslog.d/listend.conf`:

In <code>/etc/rsyslog.conf</code> :	In <code>/etc/rsyslog.d/listend.conf</code> :
<pre>\$ModLoad imuxsock \$OmitLocalLogging off</pre>	<pre>\$SystemLogSocketName ↔ /run/systemd/journal/syslog</pre>

The above lines enable rsyslog to receive information logged by journald. To enable the logging of rsyslog information in journald, we only need to add in the `/etc/rsyslog.conf`:

```
1 $Modload omjournal *.* :omjournal:
```

The above lines specify that any information being logged should be sent to `omjournal` which is a part of journald. The information from there is available from `journalctl`.

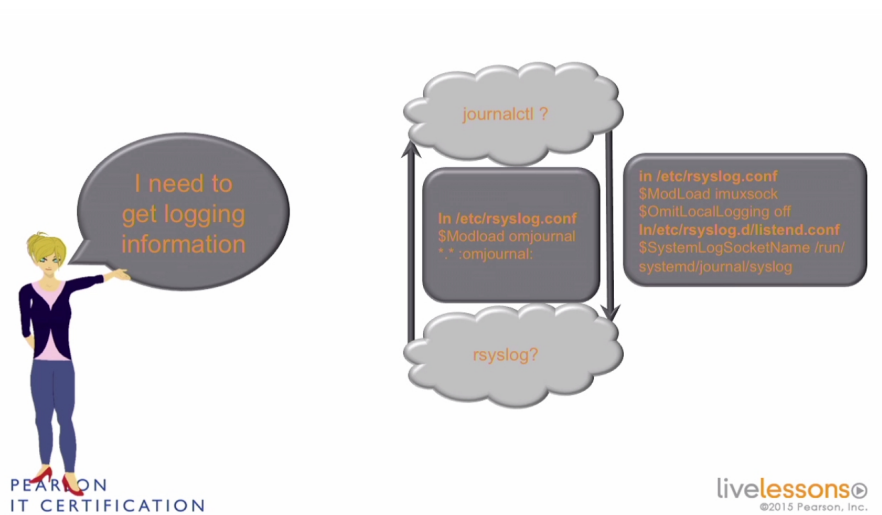


Figure 1.2: Sharing logging information between rsyslog and journald

1.2 Integrating rsyslogd and journald

Rsyslog is the old process of logging system information, and journald is the new process that's trying to do the same. Both of them are built to log process information.

1.2.1 rsyslog

Let us consider a process that's trying to write what it's doing in log files. It has two options: rsyslog and journald. If it writes to rsyslog, the advantage is that rsyslog is the system process that can handle logging for all processes. Some processes, however, just have some internal logging option (direct writes) thus bypassing rsyslog. A well known example of such a process is the *apache web server*. This is a disadvantage as if every process is doing its own logging, it's harder to manage the whole thing in a centralized way. These processes can however, be configured to write to rsyslog anyway!

1.2.2 journald

Journald is controlled by **systemd**, which takes care of starting services while booting. Systemd writes its own information to journald, and thus, if something goes wrong while starting a service, the information is available from *journald*.

For RHEL 7 in general, *journald* generally takes care of logging the startup information while *rsyslog* takes care of logging current activity of processes.

1.3 Configuring rsyslog logging

[VIDEO TUTORIAL MISSING]

1.4 Working with journald

Everything that *journald* is doing is written to a binary file. The file can be explored using two methods: by using **systemctl** and **journalctl**. Journald integrates very well with systemctl and thus, systemctl commands can get information from journald and vice versa. Using `systemctl status <serviceName>` gives us the log information that systemctl receives from the journald environment.

1.4.1 journalctl

The `journalctl` command opens up the binary file that journald is writing to. Since it's a large file, there are several filtering options:

journalctl -b

The `journalctl -b` command only shows us the boot log.

```
1 # journalctl -b
2 -- Logs begin at Sat 2017-12-02 15:25:40 IST, end at Mon 2017-12-04 13:23:04 IST. --
3 Dec 02 15:25:40 vmPrime.somuVMnet.com systemd-journal[86]: Runtime journal is using 8.0M
4 ↪ (max allowed 289.5M, trying to leave 434.3M free of 2.
5 Dec 02 15:25:40 vmPrime.somuVMnet.com kernel: Initializing cgroup subsys cpuset
6 Dec 02 15:25:40 vmPrime.somuVMnet.com kernel: Initializing cgroup subsys cpu
7 Dec 02 15:25:40 vmPrime.somuVMnet.com kernel: Initializing cgroup subsys cpuacct
8 Dec 02 15:25:40 vmPrime.somuVMnet.com kernel: Linux version 3.10.0-693.5.2.el7.x86_64
9 ↪ (builder@kbuilder.dev.centos.org) (gcc version 4.8.5 2015
10 Dec 02 15:25:40 vmPrime.somuVMnet.com kernel: Command line:
   ↪ BOOT_IMAGE=/vmlinuz-3.10.0-693.5.2.el7.x86_64 root=/dev/mapper/centos-root ro crash
11 Dec 02 15:25:40 vmPrime.somuVMnet.com kernel: Disabled fast string operations
12 ...
```

One of the best features of journald is that systemd initiates it immediately during boot and thus logs about what happens even during the very first stages of RHEL boot is available.

journalctl -since=<time>

The `journalctl` has a method to filter all results to show us what has happened since a specified period where only the logs written after a certain period are shown.

```

1 # journalctl --since=yesterday
2 -- Logs begin at Sat 2017-12-02 15:25:40 IST, end at Mon 2017-12-04 13:28:03 IST. --
3 Dec 03 23:04:36 vmPrime.somuVMnet.local systemd[1]: Time has been changed
4 Dec 03 23:04:36 vmPrime.somuVMnet.local NetworkManager[834]: <info> [1512322476.5889]
   ↪ audit: op="sleep-control" arg="off" pid=3311 uid=0 resul
5 Dec 03 23:04:36 vmPrime.somuVMnet.local systemd[1]: Stopping LSB: Bring up/down
   ↪ networking...
6 ...

```

journald -u

The `journal -u` command shows us all the logs corresponding to a certain process.

```

1 # systemctl status atd -l
2 atd.service - Job spooling tools
3 Loaded: loaded (/usr/lib/systemd/system/atd.service; enabled; vendor preset: enabled)
4 Active: active (running) since Sat 2017-12-02 15:26:18 IST; 1 day 22h ago
5 Main PID: 1224 (atd)
6 CGroup: /system.slice/atd.service
7 1224 /usr/sbin/atd -f
8
9 Dec 02 15:26:18 vmPrime.somuVMnet.local systemd[1]: Started Job spooling tools.
10 Dec 02 15:26:18 vmPrime.somuVMnet.local systemd[1]: Starting Job spooling tools...
11 # journalctl -u atd
12 -- Logs begin at Sat 2017-12-02 15:25:40 IST, end at Mon 2017-12-04 13:37:28 IST. --
13 Dec 02 15:26:18 vmPrime.somuVMnet.local systemd[1]: Started Job spooling tools.
14 Dec 02 15:26:18 vmPrime.somuVMnet.local systemd[1]: Starting Job spooling tools...

```

Both `systemctl` and `journald` are intimately interconnected. `Journald` receives it's original logging information from `systemctl`, while the information displayed by the `systemctl status` command is derived from the information stored by `journald`. To see the information in even more detail, we use the command `journalctl -u <processName> -o verbose`.

```

1 # journalctl -u atd -o verbose
2 -- Logs begin at Sat 2017-12-02 15:25:40 IST, end at Mon 2017-12-04 13:44:14 IST. --
3 Sat 2017-12-02 15:26:18.215379 IST
   ↪ [s=7e5f5839...6e;i=8f2;b=f7106...c3;m=25810b8;t=55f58...dd;x=5e9c...46]
4 PRIORITY=6
5 _UID=0
6 _GID=0
7 _BOOT_ID=f7106b6e5bc144bcac5827c5089f23c3
8 _MACHINE_ID=9d29aa554cfa4853b59f2d517a8470bd
9 SYSLOG_FACILITY=3
10 SYSLOG_IDENTIFIER=systemd
11 ...

```

The above command gives us complete information about the environment of the process.

1.5 Understanding logrotate

Logrotate is a system that's used to ensure that logging doesn't fill the hard disk of the server. Logrotate ensures that after a specified amount of time, log files will be closed and new ones opened, and a backlog of a couple of log files will be kept. This is all done

based on the assumption that old logs are useless beyond a certain age. This has a stark disadvantage of erasing logging data that may become useful at a later date.

A solution to this problem is the use of a log server, with sufficient hard disk space to keep about a year's worth of logs. The client machines can have logrotate setup to keep only a couple of week's data since there will be a backup available on the log server.

1.6 Configuring logrotate

The configuration files for logrotate reside in the `/etc` directory. The generic logrotate configuration is stored in `/etc/logrotate.conf` while the directory `/etc/logrotate.d` contains include files that RPMs dump in it for package specific log rotation. Anything in the `logrotate.d` directory will always overwrite the settings in `logrotate.conf`. Typical contents of the `logrotate.conf` file is:

```
1  # see "man logrotate" for details
2  # rotate log files weekly
3  weekly
4
5  # keep 4 weeks worth of backlogs
6  rotate 4
7
8  # create new (empty) log files after rotating old ones
9  create
10
11 # use date as a suffix of the rotated file
12 dateext
13
14 # uncomment this if you want your log files compressed
15 #compress
16
17 # RPM packages drop log rotation information into this directory
18 include /etc/logrotate.d
19
20 # no packages own wtmp and btmp -- we'll rotate them here
21 /var/log/wtmp {
22     monthly
23     create 0664 root utmp
24     minsize 1M
25     rotate 1
26 }
27
28 /var/log/btmp {
29     missingok
30     monthly
31     create 0600 root utmp
32     rotate 1
33 }
```

The last two settings demonstrate how specific logrotation instructions can be given for specific files. For example, the `/var/log/wtmp` file has to be rotated monthly, and only 1 copy of the backlog is maintained.

Since logrotate doesn't need to run all the time, it doesn't itself run as a service, but as a cron job! The config file for logrotate cron job is `/etc/cron.daily/logrotate`.

1.6.1 Checking available hard disk space

The available hard disk space and the disk space occupied by a certain directory can be checked using these two commands:

```
1 # df -h
2 Filesystem                Size      Used Avail Use% Mounted on
3 /dev/mapper/centos-root    3.8G    3.4G    412M   90% /
4 devtmpfs                  2.9G         0   2.9G    0% /dev
5 tmpfs                     2.9G         0   2.9G    0% /dev/shm
6 tmpfs                     2.9G    9.2M   2.9G    1% /run
7 tmpfs                     2.9G         0   2.9G    0% /sys/fs/cgroup
8 /dev/mapper/centos-home    7.5G     65M   7.4G    1% /home
9 /dev/mapper/centos-var     1.9G    327M   1.6G   18% /var
10 /dev/sda2                 485M    227M   258M   47% /boot
11 tmpfs                    580M     4.0K   580M    1% /run/user/42
12 tmpfs                    580M     28K   580M    1% /run/user/1000
13 /dev/sr0                  8.1G     8.1G         0 100% /run/media/somu/CentOS 7 x86_64
14 tmpfs                    580M         0   580M    0% /run/user/0
15 # du -hs /var/log
16 13M          /var/log
```
