

Chapter 1

Managing LVM Logical Volumes

1.1 Why use LVM

LVM provides a flexible approach to storage:

- Volumes can consist of more than one disk.
- Volumes can be made smaller or larger easily.
- It is easy to replace failing disks.
- Provides advanced options like working with snapshots - a method by which backups can be made of files while they're open!
- It is easy to add many new volumes. While with partitions, there is a limit of 15 partitions, there can be as many as 256 logical volumes.

1.2 Understanding LVM Setup

When working with LVM, we always start with physical storage media such as a hard disk (*sda*). Typically, a partition on the hard disk is marked as the physical volume. Now, this physical volume is added to the **volume group** which is essentially an abstraction of all the storage available. Thus, all logical volumes are created from this volume group, and from their perspective, the physical volume that's acting as it's storage media isn't important.

Once the logical volume is made from the storage in the volume group, we get a device called `/dev/vgroup/logvol`. This is the device for the logical volume on which we create the file system. As long as there is space on the volume group, we can add new logical volumes on it. If there isn't we can also add a physical volume to the volume group, to increase its capacity.




Figure 1.1: LVM Setup

1.3 Creating an LVM Logical Volume

To create a new LVM partition, first we need to make a partition like any other partition.

```

1  # fdisk /dev/sdb
2  Welcome to fdisk (util-linux 2.23.2).
3
4  Changes will remain in memory only, until you decide to write them.
5  Be careful before using the write command.
6
7
8  Command (m for help): p
9
10 Disk /dev/sdb: 10.7 GB, 10737418240 bytes, 20971520 sectors
11 Units = sectors of 1 * 512 = 512 bytes
12 Sector size (logical/physical): 512 bytes / 512 bytes
13 I/O size (minimum/optimal): 512 bytes / 512 bytes
14 Disk label type: dos
15 Disk identifier: 0xf11ab429
16
17 Device Boot      Start         End      Blocks    Id  System
18 /dev/sdb1        2048        4196351    2097152     5  Extended
19 /dev/sdb5        4096        2101247    1048576    83  Linux
20 /dev/sdb6       2103296        4196351    1046528    83  Linux
21
22 Command (m for help): n
23 Partition type:
24 p   primary (0 primary, 1 extended, 3 free)
25 l   logical (numbered from 5)
26 Select (default p): p
27 Partition number (2-4, default 2):
28 First sector (4196352-20971519, default 4196352):
29 Using default value 4196352
30 Last sector, +sectors or +size{K,M,G} (4196352-20971519, default 20971519): +100M
31 Partition 2 of type Linux and of size 100 MiB is set
32
33 Command (m for help): p
34

```

```

35 Disk /dev/sdb: 10.7 GB, 10737418240 bytes, 20971520 sectors
36 Units = sectors of 1 * 512 = 512 bytes
37 Sector size (logical/physical): 512 bytes / 512 bytes
38 I/O size (minimum/optimal): 512 bytes / 512 bytes
39 Disk label type: dos
40 Disk identifier: 0xf11ab429
41
42 Device Boot      Start          End      Blocks   Id  System
43 /dev/sdb1         2048        4196351    2097152    5  Extended
44 /dev/sdb2        4196352        4401151     102400   83   Linux
45 /dev/sdb5         4096        2101247    1048576   83   Linux
46 /dev/sdb6        2103296        4196351    1046528   83   Linux
47 # fdisk /dev/sdb
48 Welcome to fdisk (util-linux 2.23.2).
49
50 Changes will remain in memory only, until you decide to write them.
51 Be careful before using the write command.
52
53
54 Command (m for help): p
55
56 Disk /dev/sdb: 10.7 GB, 10737418240 bytes, 20971520 sectors
57 Units = sectors of 1 * 512 = 512 bytes
58 Sector size (logical/physical): 512 bytes / 512 bytes
59 I/O size (minimum/optimal): 512 bytes / 512 bytes
60 Disk label type: dos
61 Disk identifier: 0xf11ab429
62
63 Device Boot      Start          End      Blocks   Id  System
64 /dev/sdb1         2048        4196351    2097152    5  Extended
65 /dev/sdb5         4096        2101247    1048576   83   Linux
66 /dev/sdb6        2103296        4196351    1046528   83   Linux
67
68 Command (m for help): n
69 Partition type:
70 p   primary (0 primary, 1 extended, 3 free)
71 l   logical (numbered from 5)
72 Select (default p): p
73 Partition number (2-4, default 2):
74 First sector (4196352-20971519, default 4196352):
75 Using default value 4196352
76 Last sector, +sectors or +size{K,M,G} (4196352-20971519, default 20971519): +100M
77 Partition 2 of type Linux and of size 100 MiB is set
78
79 Command (m for help): p
80
81 Disk /dev/sdb: 10.7 GB, 10737418240 bytes, 20971520 sectors
82 Units = sectors of 1 * 512 = 512 bytes
83 Sector size (logical/physical): 512 bytes / 512 bytes
84 I/O size (minimum/optimal): 512 bytes / 512 bytes
85 Disk label type: dos
86 Disk identifier: 0xf11ab429
87
88 Device Boot      Start          End      Blocks   Id  System
89 /dev/sdb1         2048        4196351    2097152    5  Extended
90 /dev/sdb2        4196352        4401151     102400   83   Linux
91 /dev/sdb5         4096        2101247    1048576   83   Linux
92 /dev/sdb6        2103296        4196351    1046528   83   Linux

```

Now that we have specified the details of our new partition, we need to change one more before we can use the LVM partitions. We enter the command `t` to change the partition

type, and then show the overview of all the acceptable partition types using 1. From the command below, we can see that there is a partition type called **Linux LVM** which suits our requirements.

```

1 Command (m for help): t
2 Partition number (1,2,5,6, default 6): 2
3 Hex code (type L to list all codes): L
4
5 0 Empty          24 NEC DOS          81 Minix / old Lin bf Solaris
6 1 FAT12          27 Hidden NTFS Win 82 Linux swap / So c1 DRDOS/sec (FAT-
7 2 XENIX root     39 Plan 9            83 Linux             c4 DRDOS/sec (FAT-
8 3 XENIX usr      3c PartitionMagic  84 OS/2 hidden C:   c6 DRDOS/sec (FAT-
9 4 FAT16 <32M     40 Venix 80286       85 Linux extended   c7 Syrix
10 5 Extended       41 PPC PReP Boot    86 NTFS volume set  da Non-FS data
11 6 FAT16          42 SFS              87 NTFS volume set  db CP/M / CTOS / .
12 7 HPFS/NTFS/exFAT 4d QNX4.x            88 Linux plaintext  de Dell Utility
13 8 AIX            4e QNX4.x 2nd part  8e Linux LVM        df BootIt
14 9 AIX bootable   4f QNX4.x 3rd part  93 Amoebea          e1 DOS access
15 a OS/2 Boot Manag 50 OnTrack DM       94 Amoebea BBT       e3 DOS R/O
16 b W95 FAT32      51 OnTrack DM6 Aux 9f BSD/OS           e4 SpeedStor
17 c W95 FAT32 (LBA) 52 CP/M            a0 IBM Thinkpad hi eb BeOS fs
18 e W95 FAT16 (LBA) 53 OnTrack DM6 Aux a5 FreeBSD          ee GPT
19 f W95 Ext'd (LBA) 54 OnTrackDM6       a6 OpenBSD          ef EFI (FAT-12/16/
20 10 OPUS          55 EZ-Drive         a7 NeXTSTEP          f0 Linux/PA-RISC b
21 11 Hidden FAT12   56 Golden Bow       a8 Darwin UFS        f1 SpeedStor
22 12 Compaq diagnost 5c Priam Edisk      a9 NetBSD            f4 SpeedStor
23 14 Hidden FAT16 <3 61 SpeedStor      ab Darwin boot      f2 DOS secondary
24 16 Hidden FAT16   63 GNU HURD or Sys af HFS / HFS+        fb VMware VMFS
25 17 Hidden HPFS/NTF 64 Novell Netware  b7 BSDI fs           fc VMware VMKCORE
26 18 AST SmartSleep 65 Novell Netware  b8 BSDI swap         fd Linux raid auto
27 1b Hidden W95 FAT3 70 DiskSecure Mult bb Boot Wizard hid fe LANstep
28 1c Hidden W95 FAT3 75 PC/IX          be Solaris boot     ff BBT
29 1e Hidden W95 FAT1 80 Old Minix
30 Hex code (type L to list all codes): 8e
31 Changed type of partition 'Linux' to 'Linux LVM'
32
33 Command (m for help): p
34
35 Disk /dev/sdb: 10.7 GB, 10737418240 bytes, 20971520 sectors
36 Units = sectors of 1 * 512 = 512 bytes
37 Sector size (logical/physical): 512 bytes / 512 bytes
38 I/O size (minimum/optimal): 512 bytes / 512 bytes
39 Disk label type: dos
40 Disk identifier: 0xf11ab429
41
42 Device Boot      Start          End      Blocks   Id  System
43 /dev/sdb1        2048        4196351    2097152    5   Extended
44 /dev/sdb2        4196352    4401151      102400   8e   Linux LVM
45 /dev/sdb5        4096       2101247    1048576   83   Linux
46 /dev/sdb6       2103296    4196351    1046528   83   Linux
47
48 Command (m for help): w
49 The partition table has been altered!
50
51 Calling ioctl() to re-read partition table.
52
53 WARNING: Re-reading the partition table failed with error 16: Device or resource busy.
54 The kernel still uses the old table. The new table will be used at
55 the next reboot or after you run partprobe(8) or kpartx(8)
56 Syncing disks.
57 # partprobe

```

We enter the value 8e since it's the code for the Linux LVM partition that we need. Finally, we use p to print the partition table and verify our partition, w to save the changes and partprobe to push the changes to the kernel.

1.3.1 Creating a Physical Volume

A physical volume is just a partition with the LVM metadata added to it. The volume groups built from the PVs are not possible to build without this metadata stored in the partitions. The physical volumes are created using pvcreate. We can show all physical volumes using pvs.

```
1 # pvcreate /dev/sdb2
2 Physical volume "/dev/sdb2" successfully created.
3 # pvs
4 PV          VG      Fmt  Attr PSize   PFree
5 /dev/sda3   centos lvm2 a--  <14.91g  4.00m
6 /dev/sdb2           lvm2 ---  100.00m 100.00m
```

The pvs command tells us that we have a physical volume called /dev/sdb2 which isn't in a volume group yet, has a LVM2 formatting, has a partition size of 100MB and has the same amount of free space.

1.3.2 Creating a Volume Group

Next, we create a new volume group using the vgcreate command. Again, we can check the volume groups on our system using the vgs command.

```
1 # vgcreate vgPrime /dev/sdb2
2 Volume group "vgPrime" successfully created
3 # vgs
4 VG      #PV #LV #SN Attr   VSize   VFree
5 centos    1  4  0 wz--n- <14.91g  4.00m
6 vgPrime   1  0  0 wz--n-  96.00m 96.00m
```

The volume group *vgPrime* has 1 PV in it (/dev/sda2), no logical volumes, And has a Volume size of 96MB, all of which is free!

1.3.3 Creating a Logical Volume

The creating of a Logical Volume on a VG requires specifying the size of the logical volume. This can be done in two ways: by counting the number of extents (building blocks of LVM) [-l] or the actual size on disk (KB, MB, GB, TB, etc.)[-L]. Finally, we also can provide the name of the LV using the -n option.

```
1 # lvcreate -n lvPrime -L 96M vgPrime
2 Logical volume "lvPrime" created.
3 # lvs
4 LV      VG      Attr      LSize   Pool Origin Data%  Meta%  Move Log Cpy%Sync Convert
5 home    centos  -wi-ao---- 7.45g
6 root    centos  -wi-ao---- 3.72g
```

```
7 swap    centos  -wi-ao---- 1.86g
8 var     centos  -wi-ao---- 1.86g
9 lvPrime vgPrime -wi-a----- 96.00m
```

1.3.4 Creating a File system on the LV

Now, since the LV is ready, we can put a file system on it. We refer to the logical volume device by `/dev/<volumeGroupName>/<logicalVolumeName>`.

```
1 # mkfs.ext2 /dev/vgPrime/lvPrime
2 mke2fs 1.42.9 (28-Dec-2013)
3 Filesystem label=
4 OS type: Linux
5 Block size=1024 (log=0)
6 Fragment size=1024 (log=0)
7 Stride=0 blocks, Stripe width=0 blocks
8 24576 inodes, 98304 blocks
9 4915 blocks (5.00%) reserved for the super user
10 First data block=1
11 Maximum filesystem blocks=67371008
12 12 block groups
13 8192 blocks per group, 8192 fragments per group
14 2048 inodes per group
15 Superblock backups stored on blocks:
16 8193, 24577, 40961, 57345, 73729
17
18 Allocating group tables: done
19 Writing inode tables: done
20 Writing superblocks and filesystem accounting information: done
21
22 # mount /dev/vgPrime/lvPrime /mnt
23 # mount | grep ^/dev/
24 /dev/mapper/centos-root on / type xfs (rw,relatime,seclabel,attr2,inode64,noquota)
25 /dev/sdb5 on /data type ext4 (rw,relatime,seclabel,data=ordered)
26 /dev/mapper/centos-var on /var type xfs (rw,relatime,seclabel,attr2,inode64,noquota)
27 /dev/mapper/centos-home on /home type xfs (rw,relatime,seclabel,attr2,inode64,noquota)
28 /dev/sda2 on /boot type xfs (rw,relatime,seclabel,attr2,inode64,noquota)
29 /dev/sr0 on /run/media/somu/CentOS 7 x86_64 type iso9660
30 ↔ (ro,nosuid,nodev,relatime,uid=1000,gid=1000,iocharset=utf8,mode=0400,dmode=0500,uhelper=udisks2)
31 /dev/mapper/vgPrime-lvPrime on /mnt type ext2 (rw,relatime,seclabel)
```

Here, we can see that there is a strange behavior with the device that is mounted. While we issued the command to mount `/dev/vgPrime/lvPrime`, the device that was actually mounted is shown as `/dev/mapper/vgPrime-lvPrime`. This is because both those names are symlinks to the real name of the device (*dm-5*):

```
1 # ls -l /dev/mapper/vgPrime-lvPrime /dev/vgPrime/lvPrime
2 lrwxrwxrwx. 1 root root 7 Dec 11 15:13 /dev/mapper/vgPrime-lvPrime -> ../dm-5
3 lrwxrwxrwx. 1 root root 7 Dec 11 15:13 /dev/vgPrime/lvPrime -> ../dm-5
```

The device `/dev/dm-5` is a Device Mapper device, which is the same as used in case of LUKS encrypted volumes.

1.4 Understanding Device Mapper and LVM Device Names

The device mapper is an abstraction layer that the kernel works with to communicate with certain types of storage devices. Both LUKS encrypted partitions and LVM use the device mapper. Other devices such as software RAID and multipath also have to communicate via the device mapper.

Contrastingly, the XFS, Ext4, etc file systems work with the help of the VFS (Virtual File System) layer (instead of the Device Mapper abstraction layer). The device mapper has the devices present as `dm-*` (*dm-0, dm-1, etc.*) but we shouldn't use them. The names are assigned during boot, and are subject to change at any time! This is why the device mapper provides a bunch of symlinks to the related devices in the `/dev/mapper` directory. They are: `/dev/mapper/vg-lv` and `/dev/vg/lv` which are both symlinks to the same device.

1.5 Understanding LVM resize operations

The structure of LVMs are simple: the file system (FS) are installed on Logical Volumes (LV). These Logical Volumes get their disk space from Volume Groups (VG) which use several Physical Volumes (PV) that actually hold the data and provides the disk space.

1.5.1 Extending the File System

To expand the disk capacity of the file system, we need more space in the Logical volume. This means that (possibly) more space has to be added to the Volume Group itself, and thus, more physical volumes may need to be added.

So, first we need to create a new physical volume, and then assign it to the volume group. Then it is possible to grow the logical volume, and finally extend the file system.

1.5.2 Shrinking the File System

At first we have to reduce the size of the file system, and then reduce the size of the logical volume. If we don't there will be a file system with a bigger size than the logical partition it's residing on.

Thus, after reducing the file system size and then the logical volume size, we can then reduce the size of the volume group (if needed).

1.6 Growing an LVM Logical Volume

We can grow the LVM volume if we're running out of disk space and want to make it bigger. We typically check the amount of free space using `df -h` (*disk free - human-readable*) command:

```
1 # df -h
2 Filesystem                Size      Used Avail Use% Mounted on
3 /dev/mapper/centos-root    3.8G    3.6G   163M   96% /
4 devtmpfs                  2.9G         0   2.9G    0% /dev
5 tmpfs                     2.9G         0   2.9G    0% /dev/shm
```

```

6 tmpfs                2.9G  9.1M  2.9G   1% /run
7 tmpfs                2.9G    0  2.9G   0% /sys/fs/cgroup
8 /dev/sdb5            976M  2.6M  907M   1% /data
9 /dev/mapper/centos-var 1.9G 365M  1.5G  20% /var
10 /dev/mapper/centos-home 7.5G  68M  7.4G   1% /home
11 /dev/sda2            485M 266M  220M  55% /boot
12 tmpfs                580M  4.0K  580M   1% /run/user/42
13 tmpfs                580M  36K   580M   1% /run/user/1000
14 /dev/sr0             8.1G  8.1G    0 100% /run/media/somu/CentOS 7 x86_64
15 /dev/mapper/vgPrime-lvPrime 93M 1.6M  87M   2% /mnt

```

Now, since it's an LVM, the order in which we grow the different components matter. To make the filesystem bigger, first we check the LV size and then check to see if there's any free space in the VG:

```

1 # lvs
2 LV VG Attr LSize Pool Origin Data% Meta% Move Log Cpy%Sync
  ↔ Convert
3 root centos_cliserver -wi-ao---- <17.00g
4 swap centos_cliserver -wi-ao---- 2.00g
5 lvCLI vgCLI -wi-a----- 100.00m
6 # vgs
7 VG #PV #LV #SN Attr VSize VFree
8 centos_cliserver 1 2 0 wz--n- <19.00g 0
9 vgCLI 2 1 0 wz--n- 192.00m 92.00m

```

From the last line of the command, we can see that vgPrime has 0 VFree (i.e., free space in the VG). Thus, we need to work bottom up and make it bigger before we can make the LV and the FS bigger. So, we run `fdisk` on `/dev/sdb`.

```

1 # fdisk /dev/sdb
2 Welcome to fdisk (util-linux 2.23.2).
3
4 Changes will remain in memory only, until you decide to write them.
5 Be careful before using the write command.
6
7
8 Command (m for help): p
9
10 Disk /dev/sdb: 4294 MB, 4294967296 bytes, 8388608 sectors
11 Units = sectors of 1 * 512 = 512 bytes
12 Sector size (logical/physical): 512 bytes / 512 bytes
13 I/O size (minimum/optimal): 512 bytes / 512 bytes
14 Disk label type: dos
15 Disk identifier: 0x9287c46d
16
17 Device Boot      Start         End      Blocks    Id  System
18 /dev/sdb1          2048        206847       102400    83  Linux
19 /dev/sdb2        206848        411647       102400    83  Linux
20 /dev/sdb3        411648        821247       204800    83  Linux LVM

```

1.6.1 Creating a new logical volume in an extended partition to add to the VG

Now, we add a new partition. However, since on the given disk there's already 3 primary partitions, and there can only be a total of 4 partitions on a disk (max of 3 primary and 1

extended that can contain several logical partitions), the system defaults the last partition to be an extended one.

```
1 Command (m for help): n
2 Partition type:
3   p   primary (3 primary, 0 extended, 1 free)
4   e   extended
5 Select (default e):
6 Using default response e
7 Selected partition 4
8 First sector (821248-8388607, default 821248):
9 Using default value 821248
10 Last sector, +sectors or +size{K,M,G} (821248-8388607, default 8388607):
11 Using default value 8388607
12 Partition 4 of type Extended and of size 3.6 GiB is set
13
14 Command (m for help): p
15
16 Disk /dev/sdb: 4294 MB, 4294967296 bytes, 8388608 sectors
17 Units = sectors of 1 * 512 = 512 bytes
18 Sector size (logical/physical): 512 bytes / 512 bytes
19 I/O size (minimum/optimal): 512 bytes / 512 bytes
20 Disk label type: dos
21 Disk identifier: 0x9287c46d
22
23 Device Boot      Start          End      Blocks    Id  System
24 /dev/sdb1         2048        206847       102400    83  Linux
25 /dev/sdb2        206848        411647       102400    83  Linux
26 /dev/sdb3        411648        821247       204800    83  Linux LVM
27 /dev/sdb4        821248       8388607      3783680     5  Extended
```

Typically, we want the extended partition to take whatever disk space is left, since otherwise the space is wasted and rendered unusable due to the MBR convention used by BIOS. However, if UEFI is used, the usage of GUID (Globally Unique ID) Partition Tables (GPT) which lifts this restriction.

Now, we have to add a logical partition on the disk. Since all 4 partitions are in use, the system defaults to adding a new logical partition on the extended partition automatically. We add the new partition and then change the partition type (using `t`) to *Linux LVM* by providing the code `8e`.

```
1 Command (m for help): n
2 All primary partitions are in use
3 Adding logical partition 5
4 First sector (823296-8388607, default 823296):
5 Using default value 823296
6 Last sector, +sectors or +size{K,M,G} (823296-8388607, default 8388607): +100M
7 Partition 5 of type Linux and of size 100 MiB is set
8
9 Command (m for help): t
10 Partition number (1-5, default 5):
11 Hex code (type L to list all codes): 8e
12 Changed type of partition 'Linux' to 'Linux LVM'
13
14 Command (m for help): p
15
16 Disk /dev/sdb: 4294 MB, 4294967296 bytes, 8388608 sectors
17 Units = sectors of 1 * 512 = 512 bytes
18 Sector size (logical/physical): 512 bytes / 512 bytes
19 I/O size (minimum/optimal): 512 bytes / 512 bytes
```

```

20 Disk label type: dos
21 Disk identifier: 0x9287c46d
22
23 Device Boot      Start          End      Blocks    Id  System
24 /dev/sdb1         2048         206847       102400    83  Linux
25 /dev/sdb2        206848         411647       102400    83  Linux
26 /dev/sdb3        411648         821247       204800    83  Linux LVM
27 /dev/sdb4        821248        8388607      3783680     5  Extended
28 /dev/sdb5        823296        1028095       102400    8e  Linux LVM

```

Now we save the configuration and use the `partprobe` command to update the kernel's information about the available partitions.

```

1 Command (m for help): w
2 The partition table has been altered!
3
4 Calling ioctl() to re-read partition table.
5 Syncing disks.
6 # partprobe

```

1.6.2 Extending the Volume Group

Next let us consider we want to extend the existing LVM partition on `/dev/sdb3`. Then, we use the `vgextend` command to extend the LVM partition. It requires a Volume Group name, and a Physical device path, with the intention of adding the entire physical device to the VG. This is a *shortcut* since we don't have to create a PV on the device to be added (`/dev/sdb5`), as when all conditions are met, the `vgextend` command itself creates a PV on the disk, after which the disk is extended.

```

1 # vgextend vgCLI /dev/sdb5
2 Physical volume "/dev/sdb5" successfully created.
3 Volume group "vgCLI" successfully extended

```

Now, we can extend the logical volume to take up as much space on the VG as we want. We can confirm that our VG has been extended with the `vgs` command, and we can see which PVs are included in it (and confirm if `/dev/sdb5` is present in it), using the `pvs` command.

```

1 # vgs
2 VG                #PV #LV #SN Attr   VSize   VFree
3 centos_cliserver   1   2   0 wz--n- <19.00g     0
4 vgCLI              2   1   0 wz--n- 292.00m 192.00m
5 # pvs
6 PV                VG                Fmt  Attr PSize   PFree
7 /dev/sda2         centos_cliserver lvm2 a--  <19.00g     0
8 /dev/sdb2         lvm2  ---  100.00m 100.00m
9 /dev/sdb3         vgCLI         lvm2 a--  196.00m  96.00m
10 /dev/sdb5         vgCLI         lvm2 a--   96.00m  96.00m

```

1.6.3 Extending the LV and the File System

The LV is extended using the `lvextend` command, that takes as an argument:

Options	Description
-L	Absolute size in KiB/MiB/GiB
-l	The number of logical extents OR a percentage of either the VG size, the LV/PV size or the free space available in the VG, etc.
-r	Also resizes the file system on the LV, irrespective of file system.

The complete `lvextend` command then looks like:

```

1 # lvextend -l +100%FREE -r /dev/vgCLI/lvCLI
2 Phase 1 - find and verify superblock...
3 Phase 2 - using internal log
4 - zero log...
5 - scan filesystem freespace and inode maps...
6 - found root inode chunk
7 Phase 3 - for each AG...
8 - scan (but don't clear) agi unlinked lists...
9 - process known inodes and perform inode discovery...
10 - agno = 0
11 - agno = 1
12 - agno = 2
13 - agno = 3
14 - process newly discovered inodes...
15 Phase 4 - check for duplicate blocks...
16 - setting up duplicate extent list...
17 - check for inodes claiming duplicate blocks...
18 - agno = 0
19 - agno = 1
20 - agno = 2
21 - agno = 3
22 No modify flag set, skipping phase 5
23 Phase 6 - check inode connectivity...
24 - traversing filesystem ...
25 - traversal finished ...
26 - moving disconnected inodes to lost+found ...
27 Phase 7 - verify link counts...
28 No modify flag set, skipping filesystem flush and exiting.
29 Size of logical volume vgCLI/lvCLI changed from 100.00 MiB (25 extents) to 292.00 MiB (73
   ↔ extents).
30 Logical volume vgCLI/lvCLI successfully resized.
31 meta-data=/dev/mapper/vgCLI-lvCLI isize=512    agcount=4, agsize=6400 blks
32 =                               sectsz=512    attr=2, projid32bit=1
33 =                               crc=1          finobt=0 spinodes=0
34 data     =                               bsize=4096   blocks=25600, imaxpct=25
35 =                               sunit=0        swidth=0 blks
36 naming   =version 2                     bsize=4096   ascii-ci=0 ftype=1
37 log      =internal                       bsize=4096   blocks=855, version=2
38 =                               sectsz=512    sunit=0 blks, lazy-count=1
39 realtime =none                           extsz=4096   blocks=0, rtextents=0
40 data blocks changed from 25600 to 74752

```

The last few lines are the output from the `mkfs.xfs` command which is used to resize the file system on the disk. Had the filesystem been XFS, the `resize2fs` utility would've been used instead. The result of the operation can be verified using the `df -h` command and checking the file system size.

```

1 # df -h /dev/vgCLI/lvCLI
2 Filesystem                Size      Used Avail Use% Mounted on
3 /dev/mapper/vgCLI-lvCLI  289M    16M   274M    6% /LVM

```

1.7 Shrinking an LVM logical Volume

The shrinking operation of an LVM needs to be supported by the file system on board the LV. This is not the case for XFS as it doesn't support shrinking. **To shrink a LV, the file system on it must be unmounted first!** The size of the FS then must be reduced before shrinking the LV. To resize the Ext4 FS, we use `resize2fs` utility, which is the *xt2/Ext3/Ext4 File System Resizer*.

If we directly try to run the `resize2fs` on the disk, we'll be advised to run `e2fsck` utility to check file system consistency, i.e., if the file system has any problems with it. So, the commands to reduce the LV are:

```
1 # e2fsck -f /dev/mapper/vgCLI-lvCLI
2 e2fsck 1.42.9 (28-Dec-2013)
3 Pass 1: Checking inodes, blocks, and sizes
4 Pass 2: Checking directory structure
5 Pass 3: Checking directory connectivity
6 Pass 4: Checking reference counts
7 Pass 5: Checking group summary information
8 lvCLI: 11/25688 files (9.1% non-contiguous), 8896/102400 blocks
9 # resize2fs /dev/mapper/vgCLI-lvCLI 50M
10 resize2fs 1.42.9 (28-Dec-2013)
11 Resizing the filesystem on /dev/mapper/vgCLI-lvCLI to 51200 (1k) blocks.
12 The filesystem on /dev/mapper/vgCLI-lvCLI is now 51200 blocks long.
13 # lvreduce -L 50M /dev/mapper/vgCLI-lvCLI
14 Rounding size to boundary between physical extents: 52.00 MiB.
15 WARNING: Reducing active logical volume to 52.00 MiB.
16 THIS MAY DESTROY YOUR DATA (filesystem etc.)
17 Do you really want to reduce vgCLI/lvCLI? [y/n]: y
18 Size of logical volume vgCLI/lvCLI changed from 100.00 MiB (25 extents) to 52.00 MiB (13
   ↪ extents).
19 Logical volume vgCLI/lvCLI successfully resized.
```

Now, if there weren't any errors, we should be able to mount the file system on board the LV.

```
1 # mount /dev/mapper/vgCLI-lvCLI /LVM/
2 # mount | grep ^/dev
3 /dev/mapper/centos_cliserver-root on / type xfs
   ↪ (rw,relatime,seclabel,attr2,inode64,noquota)
4 /dev/sda1 on /boot type xfs (rw,relatime,seclabel,attr2,inode64,noquota)
5 /dev/mapper/vgCLI-lvCLI on /LVM type ext4 (rw,relatime,seclabel,data=ordered)
6 # df -h /dev/vgCLI/lvCLI
7 Filesystem              Size  Used Avail Use% Mounted on
8 /dev/mapper/vgCLI-lvCLI  45M   1.1M   40M   3% /LVM
```

In the last line we see that the file system size has been properly reduced.

1.7.1 Reduce both File system and LV in a single step

It is possible to shrink the LV and the on-board FS in a single command: (*The `-r` option automatically resizes the FS before shrinking the LV*).

```
1 # umount /LVM
2 # lvreduce -L 35M -r /dev/vgCLI/lvCLI
```

```

3  Rounding size to boundary between physical extents: 36.00 MiB.
4  fsck from util-linux 2.23.2
5  lvCLI: 11/13832 files (18.2% non-contiguous), 6886/51200 blocks
6  resize2fs 1.42.9 (28-Dec-2013)
7  Resizing the filesystem on /dev/mapper/vgCLI-lvCLI to 36864 (1k) blocks.
8  The filesystem on /dev/mapper/vgCLI-lvCLI is now 36864 blocks long.
9
10 Size of logical volume vgCLI/lvCLI changed from 52.00 MiB (13 extents) to 36.00 MiB (9
    ↪ extents).
11 Logical volume vgCLI/lvCLI successfully resized.
12 # mount /dev/mapper/vgCLI-lvCLI /LVM
13 # mount | grep ^/dev
14 /dev/mapper/centos_cliserver-root on / type xfs
    ↪ (rw,relatime,seclabel,attr2,inode64,noquota)
15 /dev/sda1 on /boot type xfs (rw,relatime,seclabel,attr2,inode64,noquota)
16 /dev/mapper/vgCLI-lvCLI on /LVM type ext4 (rw,relatime,seclabel,data=ordered)
17 # df -h /dev/mapper/vgCLI-lvCLI
18 Filesystem                Size  Used Avail Use% Mounted on
19 /dev/mapper/vgCLI-lvCLI    31M   783K   28M    3% /LVM

```

Note however, that this method won't work all the time on all file systems, due to the fact that the target FS must also support reduction via `lvreduce -r`. Thus, while the `-r` won't work on XFS, it works just fine on Ext4.